

Tutorial on Keras

CAP 6412 - ADVANCED COMPUTER VISION

SPRING 2018

KISHAN S ATHREY



Deep learning packages

- TensorFlow – Google
- PyTorch – Facebook AI research
- Keras – Francois Chollet (now at Google)
- Chainer – Company in Japan
- Caffe - Berkeley Vision and Learning Center
- CNTK - Microsoft

Python packages

Lasagne



Caffe

theano




Chainer



dmlc
mxnet



Overview of the tutorial

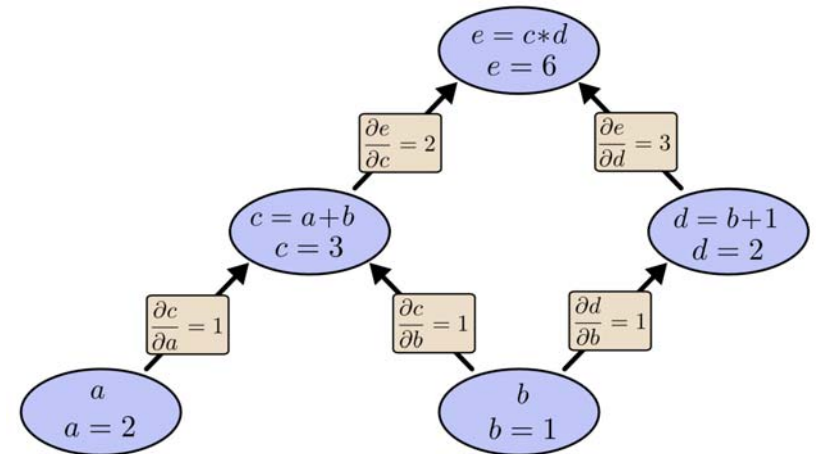
- What is Keras ?
 - Basics of Keras environment
 - Building Convolutional neural networks
 - Building Recurrent neural networks
 - Introduction to other types of layers
 - Introduction to Loss functions and Optimizers in Keras
 - Using Pre-trained models in Keras
 - Saving and loading weights and models
 - Popular architectures in Deep Learning
- 
- A solid green horizontal bar spanning the width of the slide, located at the bottom.

What is Keras ?

- **Deep neural network library in Python**
 - High-level neural networks API
 - Modular – Building model is just stacking layers and connecting computational graphs
 - Runs on top of either TensorFlow or Theano or CNTK
- **Why use Keras ?**
 - Useful for fast prototyping, ignoring the details of implementing backprop or writing optimization procedure
 - Supports Convolution, Recurrent layer and combination of both.
 - Runs seamlessly on CPU and GPU
 - Almost any architecture can be designed using this framework
 - Open Source code – Large community support


Working principle - Backend

- **Computational Graphs**
 - Expressing complex expressions as a combination of simple operations
 - Useful for calculating derivatives during backpropagation
 - Easier to implement distributed computation
 - Just specify the inputs, outputs and make sure the graph is connected



$e = c * d$
where, " $c = a + b$ " and " $d = b + 1$ "
So, $e = (a + b) * (b + 1)$
Here " a ", " b " are inputs

General pipeline for implementing an ANN


- Design and define the neural network architecture
 - Select the optimizer that performs optimization (gradient descent)
 - Select the loss function and train it
 - Select the appropriate evaluation metric for the given problem
- 
- A solid green horizontal bar spanning the width of the slide, located at the bottom.

Implementing a neural network in Keras

- **Five major steps**

- Preparing the input and specify the input dimension (size)
- Define the model architecture and build the computational graph
- Specify the optimizer and configure the learning process
- Specify the Inputs, Outputs of the computational graph (model) and the Loss function
- Train and test the model on the dataset

Note: Gradient calculations are taken care by Auto – Differentiation and parameter updates are done automatically in the backend



Procedure to implement an ANN in Keras

- Importing *Sequential class* from `keras.models`

```
from keras.models import Sequential  
  
model = Sequential()
```

- Stacking layers using `.add()` method

```
model.add(Dense(units=64, input_dim=100))  
model.add(Activation('relu'))  
model.add(Dense(units=10))  
model.add(Activation('softmax'))
```

- Configure learning process using `.compile()` method

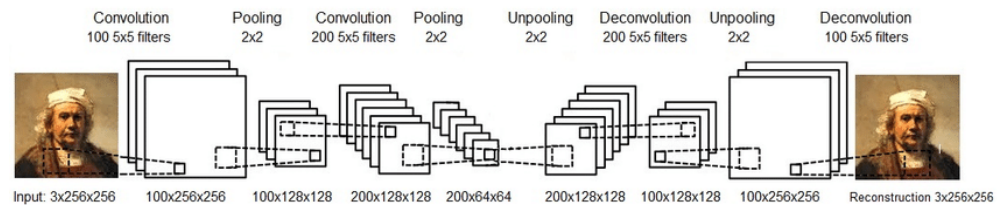
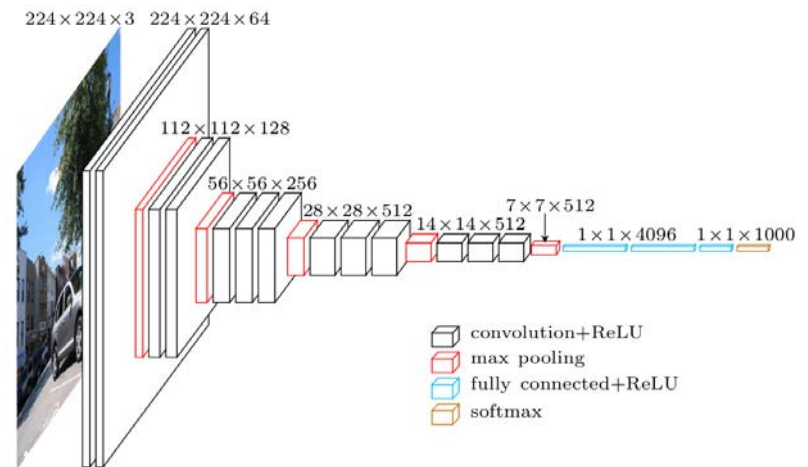
```
model.compile(loss='categorical_crossentropy',  
              optimizer='sgd',  
              metrics=['accuracy'])
```

- Train the model on train dataset using `.fit()` method

```
model.fit(x_train, y_train, epochs=5, batch_size=32)
```


Keras models – Sequential

- Sequential model
- Linear stack of layers
- Useful for building simple models
 - Simple classification network
 - Encoder – Decoder models

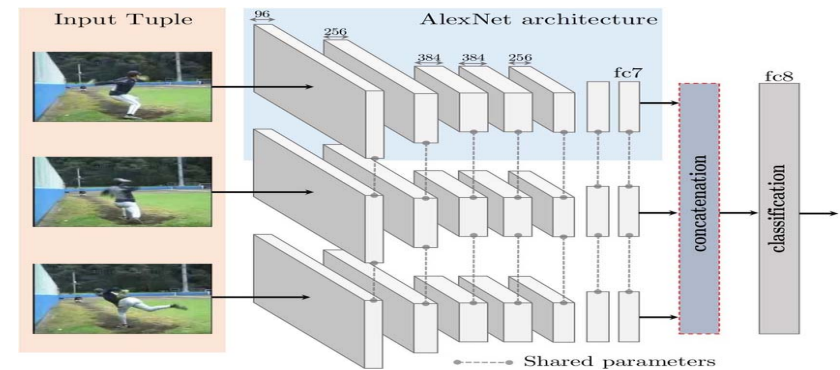


[1] <https://blog.heuritech.com/2016/02/29/a-brief-report-of-the-heuritech-deep-learning-meetup-5/vgg16/>

[2] https://www.cc.gatech.edu/~hays/7476/projects/Avery_Wenchen/

Keras models – Functional

- Functional Model
 - Multi – input and Multi – output models
 - Complex models which forks into 2 or more branches
 - Models with shared (Weights) layers



[1] <https://www.sciencedirect.com/science/article/pii/S0263224117304517>

[2] Unsupervised Domain Adaptation by Backpropagation, <https://arxiv.org/abs/1409.7495>

Keras models – Functional (Domain Adaption)



(a) MNIST

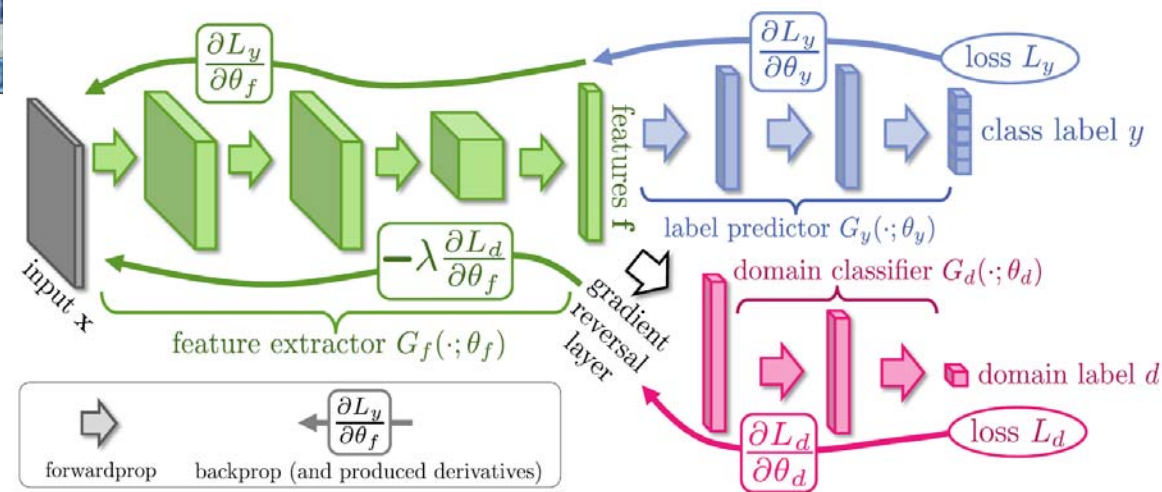
Domain A
With Labels



(b) SVHN

Domain B
Without Labels

- Train on Domain A and Test on Domain B
- Results in poor performance on test set
- The data are from different domains
- **Solution:** Adapt the model to both the domains



[1] <https://www.sciencedirect.com/science/article/pii/S0263224117304517>


[2] Unsupervised Domain Adaptation by Backpropagation, <https://arxiv.org/abs/1409.7495>


Convolution neural network - Sequential model

- Mini VGG style network
- FC – Fully Connected layers (dense layer)
- Input dimension – 4D
 - [N_Train, height, width, channels]
 - **N_train** – Number of train samples
- **Height** – height of the image
- **Width** – Width of the image
- **channels** – Number of channels
- For RGB image, **channels** = 3
- For gray scale image, **channels** = 1

```
import numpy as np
import keras
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras.optimizers import SGD

# Generate dummy data
x_train = np.random.random((100, 100, 100, 3))
y_train = keras.utils.to_categorical(np.random.randint(10, size=(100, 1)), num_classes=10)
x_test = np.random.random((20, 100, 100, 3))
y_test = keras.utils.to_categorical(np.random.randint(10, size=(20, 1)), num_classes=10)
```





```
model = Sequential()
# input: 100x100 images with 3 channels -> (100, 100, 3) tensors.
# this applies 32 convolution filters of size 3x3 each.
model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(100, 100, 3)))
model.add(Conv2D(32, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(256, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(10, activation='softmax'))

sgd = SGD(lr=0.01, decay=1e-6, momentum=0.9, nesterov=True)
model.compile(loss='categorical_crossentropy', optimizer=sgd)

model.fit(x_train, y_train, batch_size=32, epochs=10)
score = model.evaluate(x_test, y_test, batch_size=32)
```

Simple MLP network - Functional model

- Import class called “Model”
- Each layer explicitly returns a tensor
- Pass the returned tensor to the next layer as input
- Explicitly mention model inputs and outputs

```
from keras.layers import Input, Dense
from keras.models import Model

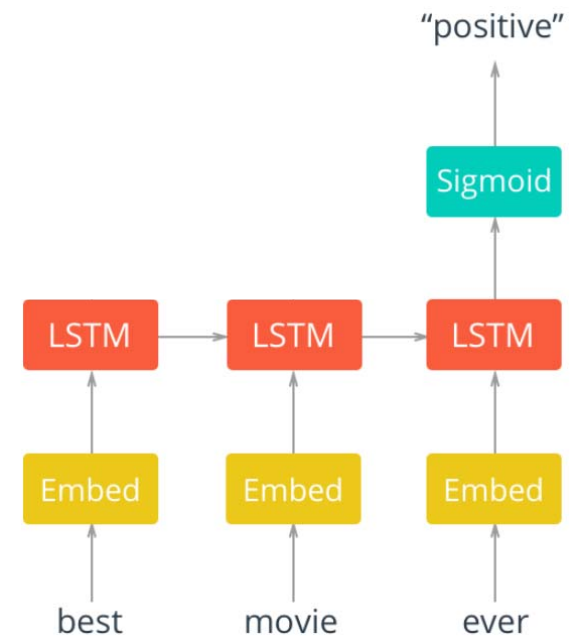
# This returns a tensor
inputs = Input(shape=(784,))

# a layer instance is callable on a tensor, and returns a tensor
x = Dense(64, activation='relu')(inputs)
x = Dense(64, activation='relu')(x)
predictions = Dense(10, activation='softmax')(x)


# This creates a model that includes
# the Input Layer and three Dense layers
model = Model(inputs=inputs, outputs=predictions)
model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
model.fit(data, labels) # starts training
```

Recurrent Neural Networks

- RNNs are used on sequential data – Text, Audio, Genomes etc.
- Recurrent networks are of three types
 - Vanilla RNN
 - LSTM
 - GRU
- They are feedforward networks with internal feedback
- The output at time “t” is dependent on current input and previous values



Recurrent Neural Network



```
from keras.models import Sequential
from keras.layers import Dense, Dropout
from keras.layers import Embedding
from keras.layers import LSTM

model = Sequential()
model.add(Embedding(max_features, output_dim=256))
model.add(LSTM(128))
model.add(Dense(1, activation='sigmoid'))

model.compile(loss='binary_crossentropy',
              optimizer='rmsprop',
              metrics=['accuracy'])

model.fit(x_train, y_train, batch_size=16, epochs=10)
score = model.evaluate(x_test, y_test, batch_size=16)
```


Convolution layers

- 1D Conv

`keras.layers.convolutional.Conv1D(filters, kernel_size, strides=1, padding='valid', dilation_rate=1, activation=None, use_bias=True, kernel_initializer='glorot_uniform', bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None, kernel_constraint=None, bias_constraint=None)`

Applications: Audio signal processing, Natural language processing

- 2D Conv

`keras.layers.convolutional.Conv2D(filters, kernel_size, strides=(1, 1), padding='valid', data_format=None, dilation_rate=(1, 1), activation=None, use_bias=True, kernel_initializer='glorot_uniform', bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None, kernel_constraint=None, bias_constraint=None)`

Applications: Computer vision - Images

- 3D Conv


`keras.layers.convolutional.Conv3D(filters, kernel_size, strides=(1, 1, 1), padding='valid', data_format=None, dilation_rate=(1, 1, 1), activation=None, use_bias=True, kernel_initializer='glorot_uniform', bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None, kernel_constraint=None, bias_constraint=None)`

Applications: Computer vision – **Videos** (Convolution along temporal dimension)

Pooling layers

- Max pool

`keras.layers.pooling.MaxPooling2D(pool_size=(2, 2), strides=None, padding='valid')`




- Average pool

`keras.layers.pooling.AveragePooling2D(pool_size=(2, 2), strides=None, padding='valid')`

- Up sampling

`keras.layers.convolutional.UpSampling2D(size=(2, 2))`



General layers

- Dense

keras.layers.core.Dense(units, activation=**None**, use_bias=**True**, kernel_initializer='glorot_uniform', bias_initializer='zeros', kernel_regularizer=**None**, bias_regularizer=**None**, activity_regularizer=**None**, kernel_constraint=**None**, bias_constraint=**None**)

- Dropout

keras.layers.core.Dropout(rate, noise_shape=**None**, seed=**None**)

- Embedding

keras.layers.embeddings.Embedding(input_dim, output_dim, input_length=**None**, embeddings_initializer='uniform', embeddings_regularizer=**None**, activity_regularizer=**None**, embeddings_constraint=**None**, mask_zero=**False**)

Optimizers available in Keras

- How do we find the “best set of parameters (weights and biases)” for the given network ?
- **Optimization**
 - They vary in the speed of convergence, ability to avoid getting stuck in local minima
 - SGD – Stochastic gradient descent
 - SGD with momentum
 - Adam
 - AdaGrad
 - RMSprop
 - AdaDelta
- Detailed explanation of each optimizer is given in the “Deep learning book”
 - URL: <http://www.deeplearningbook.org/contents/optimization.html>

Loss functions available in Keras

- MSE – Mean square error

$$\text{MSE} = \frac{1}{n} \sum_{t=1}^n e_t^2$$

- MAE – Mean absolute error

$$\text{MAE} = \frac{1}{n} \sum_{t=1}^n |e_t|$$

- Categorical cross entropy – “K” number of classes

$$J(\theta) = - \left[\sum_{i=1}^m \sum_{k=1}^K 1 \{y^{(i)} = k\} \log P(y^{(i)} = k | x^{(i)}; \theta) \right]$$

- KL divergence – If $P(X)$ and $Q(X)$ are two different probability distributions, then we can measure how different these two distributions are using KL divergence

$$D_{\text{KL}}(P \| Q) = \mathbb{E}_{x \sim P} \left[\log \frac{P(x)}{Q(x)} \right] = \mathbb{E}_{x \sim P} [\log P(x) - \log Q(x)]$$

Loading and Saving Keras models

- Use *.save* method to save the model
- Use *load_model* function to load saved model
- Saved file contains –
 - Architecture of the model
 - Weights and biases
 - State of the optimizer
- Saving weights
- Loading all the weights and loading weights layer wise

```
from keras.models import load_model

model.save('my_model.h5') # creates a HDF5 file 'my_model.h5'
del model # deletes the existing model

# returns a compiled model
# identical to the previous one
model = load_model('my_model.h5')
```

```
model.save_weights('my_model_weights.h5')
model.load_weights('my_model_weights.h5', by_name=True)
```

Extracting features from pre-trained models

- Import the network [eg:VGG16]
- Specify the weights
- Specify whether the classifier at the top has to be included or not
- The argument *“include_top = False”* – removes the classifier from the imported model
- The input size of the image must be same as what the imported model was trained on (with exceptions)

```
from keras.applications.vgg16 import VGG16
from keras.preprocessing import image
from keras.applications.vgg16 import preprocess_input
import numpy as np

model = VGG16(weights='imagenet', include_top=False)

img_path = 'elephant.jpg'
img = image.load_img(img_path, target_size=(224, 224))
x = image.img_to_array(img)
x = np.expand_dims(x, axis=0)
x = preprocess_input(x)


features = model.predict(x)
```

Popular Deep learning Architectures

- Popular Convolution networks
 - Alex net
 - VGG
 - Res-Net
 - DenseNet
- Generative models
 - Autoencoders
 - Generative adversarial networks

Image recognition networks

- AlexNet – 2012



- VGG - 2014

VGGNet




[1] AlexNet, <https://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks.pdf>


[2] VGG Net, <https://arxiv.org/pdf/1409.1556.pdf>

Image recognition networks

- ResNet – 2015 (residual connections)



- DenseNet – 2017 (Dense connectivity)




[1] ResNet, <https://arxiv.org/abs/1512.03385>


[2] DenseNet, <https://arxiv.org/abs/1608.06993>

Performance of the recognition networks

Revolution of Depth



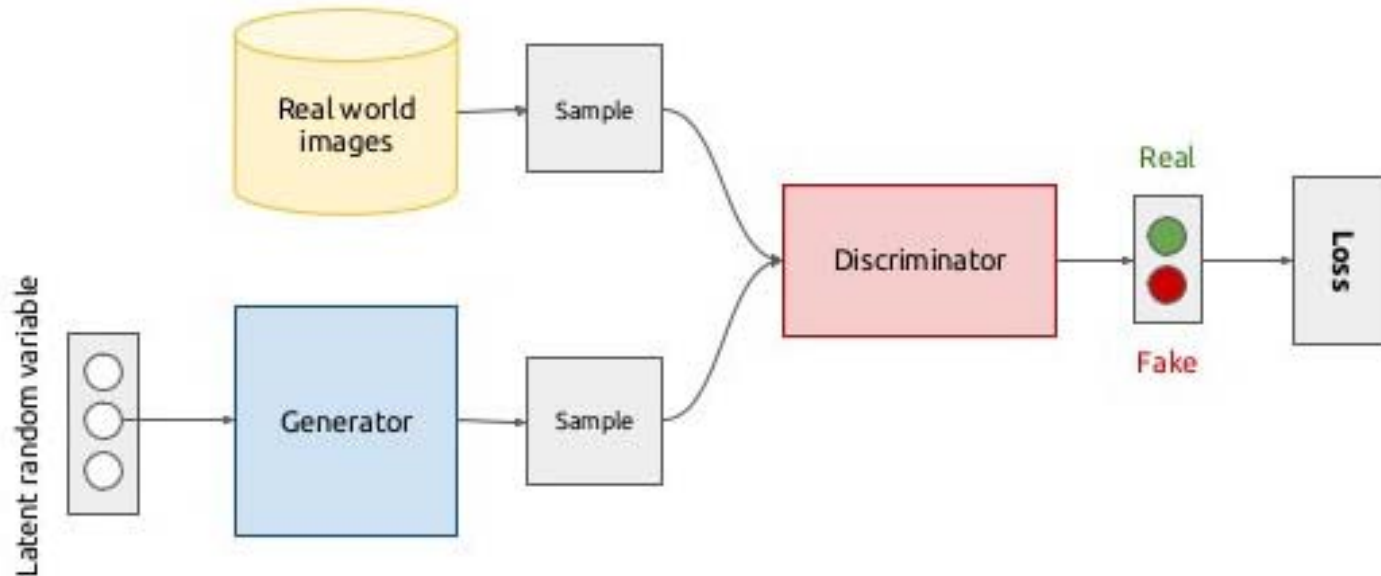
Autoencoders



- Unsupervised representation learning
- Dimensionality reduction
- Denoising




Generative Adversarial Network



Interesting Applications using GANs

- Generate images from textual description
- Performing arithmetic in latent space



This small blue bird has a short pointy beak and brown on its wings



This bird is completely red with black wings and pointy beak



A small sized bird that has a cream belly and a short pointed bill



A small bird with a black head and wings and features grey wings



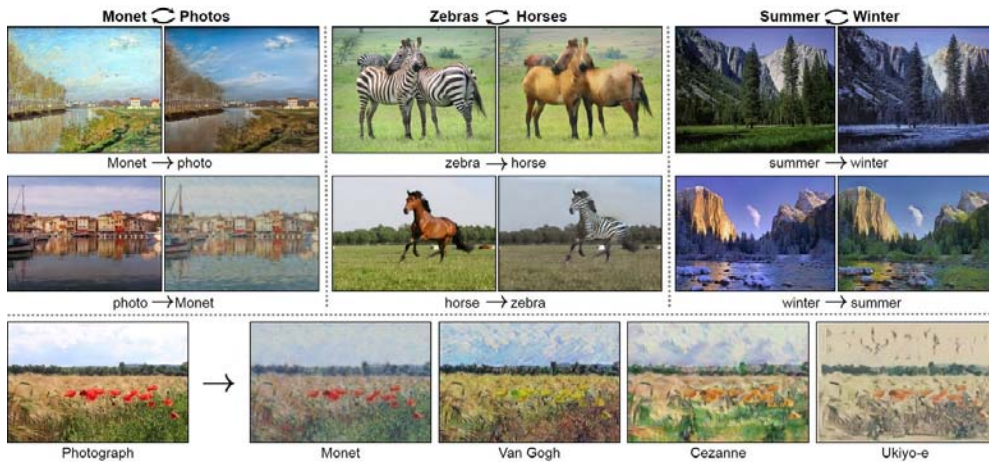
[1] Stack GAN, <https://arxiv.org/abs/1612.03242>

[2] DC GAN, <https://arxiv.org/abs/1511.06434>

Interesting Applications using GANs



- Generate images of the same scene with different weather conditions
- Transfer the style of painting from one image to other
- Change the content in the image



[1] UNIT, <https://arxiv.org/pdf/1703.00848>

[2] Cyclic GAN, <https://arxiv.org/abs/1703.10593>

Community contributed layers and other functionalities

https://github.com/farizrahman4u/keras-contrib/tree/master/keras_contrib

<https://github.com/fchollet/keras/tree/master/keras/layers>

Keras Documentation – keras.io

Keras Blog - <https://blog.keras.io/index.html>

Questions ?
